

Excellent.

Now we move into **PHASE 2**, which is where many engineers *think* they are strong—but most are not.

I'll explain **Phase 2 completely, cleanly, and practically**, exactly how it must be done on a real project, without skipping reasoning steps or jumping to registers.

PHASE 2 – CORE ARCHITECTURE & PROCESSING MODEL

(How computation actually happens inside the chip)

Duration

 1–2 weeks (depending on background)

Phase-2 Goal

Build a **mental execution model** of the processor
so that you can predict behavior, debug failures, and reason about performance.

This phase answers:

- *What happens from reset to first instruction?*
 - *How does a core execute code?*
 - *How does memory access really work?*
 - *Why multicore is hard?*
-

1 What Phase 2 IS and IS NOT

Phase 2 IS:

- Understanding **execution flow**
- Understanding **core, cache, memory relationships**
- Understanding **privilege and control**
- Understanding **multi-core interactions**

Phase 2 is NOT:

- ✗ Writing drivers
- ✗ Tuning performance

✗ Studying every register

✗ OS internals

2 Step 1: Reset → First Instruction (The Most Important Flow)

If you understand this, **everything else becomes logical**.

Questions you must answer

- Who applies reset?
 - What logic is alive after reset?
 - Where is the first instruction fetched from?
 - Who decides the boot source?
-

General processor flow

```
Power stable
→ Reset deasserted
→ Core released
→ Reset vector read
→ First instruction fetched
```

For T2080 (Power Architecture)

- Reset releases **one core**
- Core fetches instruction from **Boot ROM**
- Boot ROM behavior is influenced by **RCW**
- Everything else is disabled initially

★ Key insight:

Nothing runs magically. Hardware decides first.

3 Step 2: CPU Core Basics (Without Micro-Details)

You do NOT start with pipeline stages.

What you must understand:

- Instruction fetch
- Decode

- Execute
- Write-back (conceptually)

Power Architecture essentials

- RISC-based
- Load/store architecture
- Fixed instruction width (mostly)
- Strong deterministic behavior

❖ Why this matters:

- Predictable execution
- Easier real-time guarantees
- Different from ARM/x86 assumptions

4 Step 3: Privilege Levels & Control

This is critical for debugging and OS interaction.

Typical privilege levels:

- User mode
- Supervisor / Kernel mode
- Hypervisor (in advanced SoCs)

Key questions:

- Who can touch hardware registers?
- Who configures MMU?
- Who controls cache?

For T2080:

- Early boot runs in **highest privilege**
- OS later drops privilege
- Some registers become locked

❖ Debug insight:

If you don't know privilege, you'll chase "permission bugs" forever.

5 Step 4: Memory Hierarchy (Core → Cache → DDR)

This is where most real bugs hide.

Memory layers:

```
Core registers
  → L1 cache
  → L2 cache
  → (L3 cache, if present)
  → DDR
```

What you must understand:

- Why cache exists
- Latency vs bandwidth
- Cache line concept
- Cache coherency

📌 Golden rule:

Memory is slow. Cache hides that fact.

6 Step 5: Cache Coherency (Multicore Reality)

Multicore without coherency = chaos.

You must understand:

- What cache coherency means
- Why hardware fabric exists
- What happens when coherency breaks

Real-world failure:

- Random crashes
 - Stale data
 - Non-reproducible bugs
- 📌 This is why networking SoCs invest heavily in coherency fabric.
-

7 Step 6: Endianness (Silent Killer)

Endianness errors look like “random data corruption”.

Concepts:

- Big-endian vs little-endian
- How memory is laid out
- Why networking prefers big-endian

T2080 context:

- Power Architecture traditionally big-endian
- Software must agree with hardware

📌 Debug signal:

Correct values in registers, wrong in memory = endianness mismatch.

8 Step 7: MMU Basics (Without OS Deep Dive)

You must know:

- Virtual vs physical memory
- Address translation concept
- Why MMU exists

Why this matters:

- Device memory vs normal memory
- Cacheable vs non-cacheable
- DMA correctness

📌 Bring-up insight:

Wrong MMU settings cause invisible bugs.

9 Step 8: Multicore Boot Model

Key questions:

- Which core boots first?
- How are other cores released?
- Who initializes shared resources?

Typical model:

- Core 0 boots
- Initializes memory, cache
- Releases other cores

Bug pattern:

Two cores touching hardware too early = deadlock.

Step 9: Phase-2 Practical Exercises (Mandatory)

Exercise 1: Instruction Flow Diagram

Draw reset → first instruction → memory access path.

Exercise 2: Cache Thought Experiment

Explain what happens when two cores write the same variable.

Exercise 3: Memory Access Path

Trace one load instruction from core to DDR.

Step 10: Phase-2 Output (Non-Negotiable)

Before moving forward, the junior must produce:

A clean explanation (verbal or written) of:

1. Reset-to-execution flow
2. Core → cache → memory path
3. Cache coherency importance
4. Privilege & control roles
5. Multicore startup sequence

Evaluation criteria:

- No jargon dumping
 - Correct sequencing
 - Real-world awareness
-

Phase-2 Failure Signals

Phase 2 has failed if the junior:

- ✗ Talks about cache as “speed booster” only
- ✗ Cannot explain first instruction fetch
- ✗ Confuses physical and virtual memory

- ✖ Ignores coherency
 - ✖ Cannot explain multicore behavior
-

⌚ Phase-2 Completion Statement

When Phase 2 is done, the junior should be able to say:

"I can mentally simulate how code executes on this processor, even before running it."

That's the **real milestone**.

⌚ What Comes Next

Phase 3 will bring this knowledge into **real hardware responsibility**:

- Power domains
- Reset trees
- Clocking
- Sequencing failures

This is where boards live or die.

If you want, next I can:

- Proceed to **Phase 3 – Power, Reset & Clock**
- Convert Phase 2 into a **training worksheet**
- Show **Phase-2 interview questions with ideal answers**

Just tell me 